

Diego Atencia

Data Engineering — SQL · Python · AWS

Madrid, Spain • dratencia@gmail.com • +34 644 606 287
github.com/alektebel • alektebel.github.io

SUMMARY

Mathematics graduate (UAM) with 3+ years working on banking data: regulatory reporting pipelines (FINREP, COREP, CIRBE) on Oracle and T-SQL, credit-risk (LGD) modeling, and a self-built full-stack application deployed on AWS. Expert SQL, strong Python. Based in Madrid.

EXPERIENCE

Senior Data Scientist

Jun 2025 – Present

PwC — Madrid

- Data engineering and analysis on the LGD database of a top-4 Spanish bank, where small errors have significant impact on the bank's reserved capital and earnings.
- Built an AWS application with an Angular frontend to streamline and automate the production of audit reporting documents.
- Developing a data-quality-control (DQC) automation tool, currently at proof-of-concept stage.
- Contributed to landing 8 contract renewals with this client; integrated 2 junior engineers now working full-time on the account.

Data Scientist

Nov 2023 – Jun 2025

PwC — Madrid

- Worked on stress testing for a PD (Probability of Default) climate-risk application in Excel.

SQL & PL/SQL Developer

Oct 2022 – Nov 2023

Álamo Consulting — Madrid

- Owned regulatory reporting pipelines (FINREP, COREP, SIRBE/CIRBE) for an external banking entity using Oracle and T-SQL.
- Built scalable transformations within a large regulatory codebase, ensuring accurate and timely data delivery.

EDUCATION

B.Sc. in Mathematics

Sep 2018 – Jun 2023

Autonomous University of Madrid (UAM)

Final thesis: Weak solutions to Navier–Stokes equations (9/10). Founder of university AI student association (70+ members).

M.Sc. in Artificial Intelligence (8.67/10)

Sep 2024 – Oct 2025

International University of Valencia

Thesis: *Optimizing low-cost hand prosthesis control using Deep Reinforcement Learning.*

OPEN-SOURCE CONTRIBUTIONS

Liger-Kernel — github.com/linkedln/Liger-Kernel

PR #1128 (under review)

Production CUDA kernel library for LLM training, used at scale across industry

- Debugging a distributed-training bug in **LigerTiledSwiGLUMLP**: FSDP's post-backward reshard hooks were triggered prematurely inside a tiling loop, corrupting parameter shards across GPUs.

LEANN

Issue #41

- Proposed a hybrid BM25+RRF retrieval pipeline over the neural-only architecture to improve recall on sparse queries.

SELECTED PROJECT

Banking-Regulation LLM — fine-tuned model for Spanish banking regulation Q&A

2025–2026

- Fine-tuned Qwen2.5-7B with LoRA (model published on Hugging Face); built the RAG pipeline and AWS deployment (ECS Fargate, RDS/pgvector, ALB) around it; served real queries during early access.

TECHNICAL SKILLS

Languages	SQL (expert: Oracle PL/SQL, T-SQL), Python (expert), C/C++ (strong)
Systems & Cloud	AWS, Docker, Linux, Git/GitHub, CI/CD principles, distributed training (FSDP/DDP), CUDA kernels
Databases	Oracle, PostgreSQL, ChromaDB (vector), SAS
ML / GenAI	PyTorch, TensorFlow, Scikit-learn, HuggingFace, LoRA/QLoRA, RAG pipelines
CS Fundamentals	Data structures & algorithms, OOP, distributed systems, debugging complex systems, code review

LANGUAGES

Spanish (native)	English (fluent, professional working proficiency)
------------------	--